

Variable Importance in Machine Learning

1. Introduction and Intuition

In machine learning, we are often interested not only in **predictive performance**, but also in understanding **which variables (features) drive those predictions**. This concept is known as *variable importance*.

Understanding variable importance helps us:

- Interpret model behavior
- Identify key drivers of outcomes
- Perform feature selection
- Detect data issues (e.g., leakage, redundancy)
- Generate domain insights (e.g., in healthcare, economics)

A critical guiding principle:

Variable importance is only meaningful if the model itself performs well.

2. General Setup and Notation

We assume a dataset:

$$\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$$

where:

- $x_i = (x_{i1}, x_{i2}, \dots, x_{ip})$ is a feature vector
- y_i is the response
- $\hat{f}(x)$ is a trained model

We aim to quantify how much each feature x_j contributes to predictive performance.

3. Linear Regression

Model

$$\hat{y} = \beta_0 + \sum_{j=1}^p \beta_j x_j$$

Intuition

Each coefficient β_j represents the change in prediction for a one-unit increase in x_j , holding other variables constant.

Example: House Price Prediction

$$\hat{y} = 50,000 + 150 \cdot x_1 + 20,000 \cdot x_2 - 5,000 \cdot x_3$$

- x_1 : square footage
- x_2 : number of bedrooms
- x_3 : age of house

Interpretation:

- Larger homes -> higher prices
- More bedrooms -> higher prices
- Older homes -> lower prices

Numerical comparison (raw scale):

- $|\beta_1| = 150, |\beta_2| = 20,000, |\beta_3| = 5,000$ -> Bedrooms appears most important

After standardization (example results):

- $|\beta_1^{(\text{std})}| = 0.80$
 - $|\beta_2^{(\text{std})}| = 0.40$
 - $|\beta_3^{(\text{std})}| = 0.20$ -> Square footage is actually most important
-

Measuring Importance

Naively:

$$\text{Importance}(x_j) = |\beta_j|$$

Better (after standardization):

$$\text{Importance}(x_j) = |\beta_j^{(\text{std})}|$$

This reflects the effect of a **1 standard deviation change**.

Pitfalls

1. **Scale dependence** Coefficients are not comparable unless features are standardized.
2. **Multicollinearity** Correlated features split importance and produce unstable coefficients.

3. **Model misspecification** If the true relationship is nonlinear, coefficients may mislead.
 4. **Omitted variable bias** Missing variables distort importance of included ones.
 5. **Poor model performance** If R^2 is low, importance is unreliable.
-

4. Logistic Regression

Model

$$\hat{p}(x) = \frac{1}{1 + e^{-(\beta_0 + \sum_{j=1}^p \beta_j x_j)}}$$

Intuition

Coefficients describe changes in **log-odds**:

$$\beta_j = \frac{\partial}{\partial x_j} \log\left(\frac{p}{1-p}\right)$$

Example: Disease Risk Prediction

$$\log\left(\frac{p}{1-p}\right) = -5 + 0.04x_1 + 1.2x_2 + 0.8x_3$$

- x_1 : age
- x_2 : smoker (0/1)
- x_3 : high blood pressure (0/1)

Interpretation:

- Smoking increases odds of disease significantly

Numerical comparison (odds ratios):

- Age: $e^{0.04} \approx 1.04$ (4% increase per year)
- Smoking: $e^{1.2} \approx 3.32$
- Blood pressure: $e^{0.8} \approx 2.23$

-> Smoking has the strongest effect

- Odds ratio: $e^{1.2} \approx 3.32$
-

Measuring Importance

- Standardized coefficients
 - Odds ratios e^{β_j}
-

Pitfalls

1. **Scale issues** Continuous vs binary features not directly comparable.
 2. **Log-odds interpretation** Coefficients are not direct probability changes.
 3. **Class imbalance** Rare events can lead to unstable estimates.
 4. **Correlated predictors** Same instability issues as linear regression.
 5. **Poor model performance** Low AUC or accuracy -> unreliable importance.
-

5. Tree-Based Models (Decision Trees and Random Forests)

Intuition

Tree models split data to reduce impurity. A feature is important if it produces large reductions in impurity.

Impurity Measures

Classification (Gini):

$$G(t) = 1 - \sum_{k=1}^K p_k^2$$

Regression (Variance):

$$\text{Var}(t) = \frac{1}{|t|} \sum_{i \in t} (y_i - \bar{y}_t)^2$$

Importance Formula

$$\text{Importance}(x_j) = \sum_{t \in \text{splits on } x_j} \Delta I(t)$$

where:

$$\Delta I(t) = I(\text{parent}) - \left(\frac{n_L}{n} I(\text{left}) + \frac{n_R}{n} I(\text{right}) \right)$$

Example: Loan Default Prediction

Features:

- income

- credit score
- zip code
- number of late payments

Typical behavior:

- Early splits on credit score -> large impurity reduction
- Later splits on zip code -> small reduction

Numerical comparison (example importance scores):

- Credit score: 0.45
- Income: 0.30
- Late payments: 0.20
- Zip code: 0.05

Conclusion:

- Credit score is highly important
 - Zip code is less important
-

Pitfalls

1. **Bias toward high-cardinality features** Variables with many unique values may appear overly important.
 2. **Correlated features** Importance is shared across correlated variables.
 3. **Model instability (single trees)** Small changes in data -> different splits.
 4. **Early split bias** Features used near the root dominate importance.
 5. **Poor model performance** If the tree predicts poorly, importance is unreliable.
-

6. Permutation Importance (Model-Agnostic)

Intuition

If a feature is important, breaking its relationship with the response should degrade model performance.

Procedure

1. Compute baseline loss:

$$L_{\text{orig}} = \frac{1}{n} \sum_{i=1}^n \ell(y_i, \hat{f}(x_i))$$

2. Shuffle feature x_j .
3. Compute new loss:

$$L_{\text{perm}(j)}$$

4. Importance:

$$\text{Importance}(x_j) = L_{\text{perm}(j)} - L_{\text{orig}}$$

Example: Student Performance Prediction

Features:

- study hours
- attendance
- parental education
- favorite color

Results:

- Permuting study hours -> large drop in accuracy
- Permuting favorite color -> no change

Numerical comparison (accuracy drops):

- Study hours: 0.85 -> 0.65 (drop = 0.20)
- Attendance: 0.85 -> 0.75 (drop = 0.10)
- Parental education: 0.85 -> 0.82 (drop = 0.03)
- Favorite color: 0.85 -> 0.85 (drop = 0.00)

Conclusion:

- Study hours is important
 - Favorite color is irrelevant
-

Pitfalls

1. **Correlated features** Importance may be underestimated if correlated variables remain.
 2. **Unrealistic data combinations** Permutation can break natural relationships in data.
 3. **Computational cost** Requires repeated model evaluation.
 4. **Variance in estimates** Results may vary across permutations.
 5. **Poor model performance** If baseline model is weak, all features may appear unimportant.
-

7. Model Complexity vs Interpretability

An important consideration when working with variable importance is the tradeoff between **model complexity** and **interpretability**.

Intuition

- **Simple models** (e.g., linear and logistic regression):
 - Easy to interpret
 - Variable importance is directly tied to coefficients
 - May fail to capture complex relationships
 - **Complex models** (e.g., random forests, gradient boosting, neural networks):
 - Capture nonlinearities and interactions
 - Often achieve better predictive performance
 - Harder to interpret
-

Implications for Variable Importance

1. **Interpretability vs Accuracy Tradeoff** A highly interpretable model may provide clear importance measures but lower predictive accuracy, while a complex model may be more accurate but harder to explain.
 2. **Different Notions of Importance**
 - Linear models: importance is *explicit and local* (coefficients)
 - Tree-based and black-box models: importance is *implicit and global* (performance-based or impurity-based)
 3. **Use of Model-Agnostic Methods** Techniques like permutation importance help bridge this gap by providing interpretability for complex models.
 4. **Risk of Misleading Simplicity** A simple model may give clean interpretations that are actually incorrect if the model is misspecified.
-

Practical Guidance

- If **interpretability is critical** (e.g., healthcare, policy):
 - Prefer simpler models or combine with interpretability tools
- If **prediction is the main goal**:
 - Use more complex models
 - Apply model-agnostic importance methods
- In practice, it is often useful to:

- Fit both simple and complex models
- Compare variable importance across them

8. Side-by-Side Comparison on a Single Dataset

To see how variable importance depends on the model, consider one shared dataset: a **loan default prediction** problem.

Dataset

- **Outcome:** default within 12 months (yes/no)
- **Features:**
 - credit score
 - debt-to-income ratio
 - number of late payments
 - annual income

This dataset is useful because the same features can be analyzed with several modeling approaches.

Results Across Methods

Method	Model performance	Most important feature(s)	Numerical importance summary
Linear regression (linear probability model, for comparison)	Test AUC = 0.77	Late payments, credit score	Standardized coefficients: late payments ($\beta = 0.42$), credit score ($\beta = 0.38$), debt-to-income ($\beta = 0.25$), income ($\beta = 0.10$)
Logistic regression	Test AUC = 0.84	Credit score, late payments	Odds ratios: credit score ($e^{-0.85} = 0.43$), late payments ($e^{1.10} = 3.00$), debt-to-income ($e^{0.55} = 1.73$), income ($e^{-0.20} = 0.82$)
Random forest (impurity importance)	Test accuracy = 0.86	Credit score, late payments	Gini-based importance: credit score 0.34, late payments 0.31, debt-to-income 0.21, income 0.14

Random forest (permutation importance)	Test accuracy = 0.86	Late payments, credit score	Accuracy drop after permutation: late payments 0.11, credit score 0.09, debt-to- income 0.05, income 0.02
--	-------------------------	-----------------------------------	---

What This Shows

- The **same dataset** can produce different rankings depending on the model and importance definition.
- The tree-based methods and permutation importance often agree on the strongest predictors, but not always in the exact same order.
- The linear and logistic models provide coefficient-based interpretations, while the random forest uses split-based and performance-based notions of importance.
- Because the linear probability model has weaker classification performance than logistic regression and random forest, its importance summary is less trustworthy.

Interpretation

In this example, **credit score** and **late payments** are consistently among the most important variables. **Income** is comparatively less important across all methods. The exact ranking shifts because each method measures importance differently:

- coefficients measure directional effect size,
- tree impurity importance measures split quality,
- permutation importance measures predictive dependence.

9. Key Takeaways

- Variable importance depends on the **model and data**.
- Different methods can yield **different rankings**.
- Correlation among features complicates interpretation.
- No importance measure implies causality.

Most Important Principle

Always evaluate model performance before interpreting variable importance.

10. Practical Recommendations

- Standardize features for linear models
- Use multiple importance methods when possible
- Prefer permutation importance for model-agnostic insight
- Check for correlated predictors
- Validate results with domain knowledge

This framework provides a foundation for understanding, computing, and critically evaluating variable importance across common machine learning models.