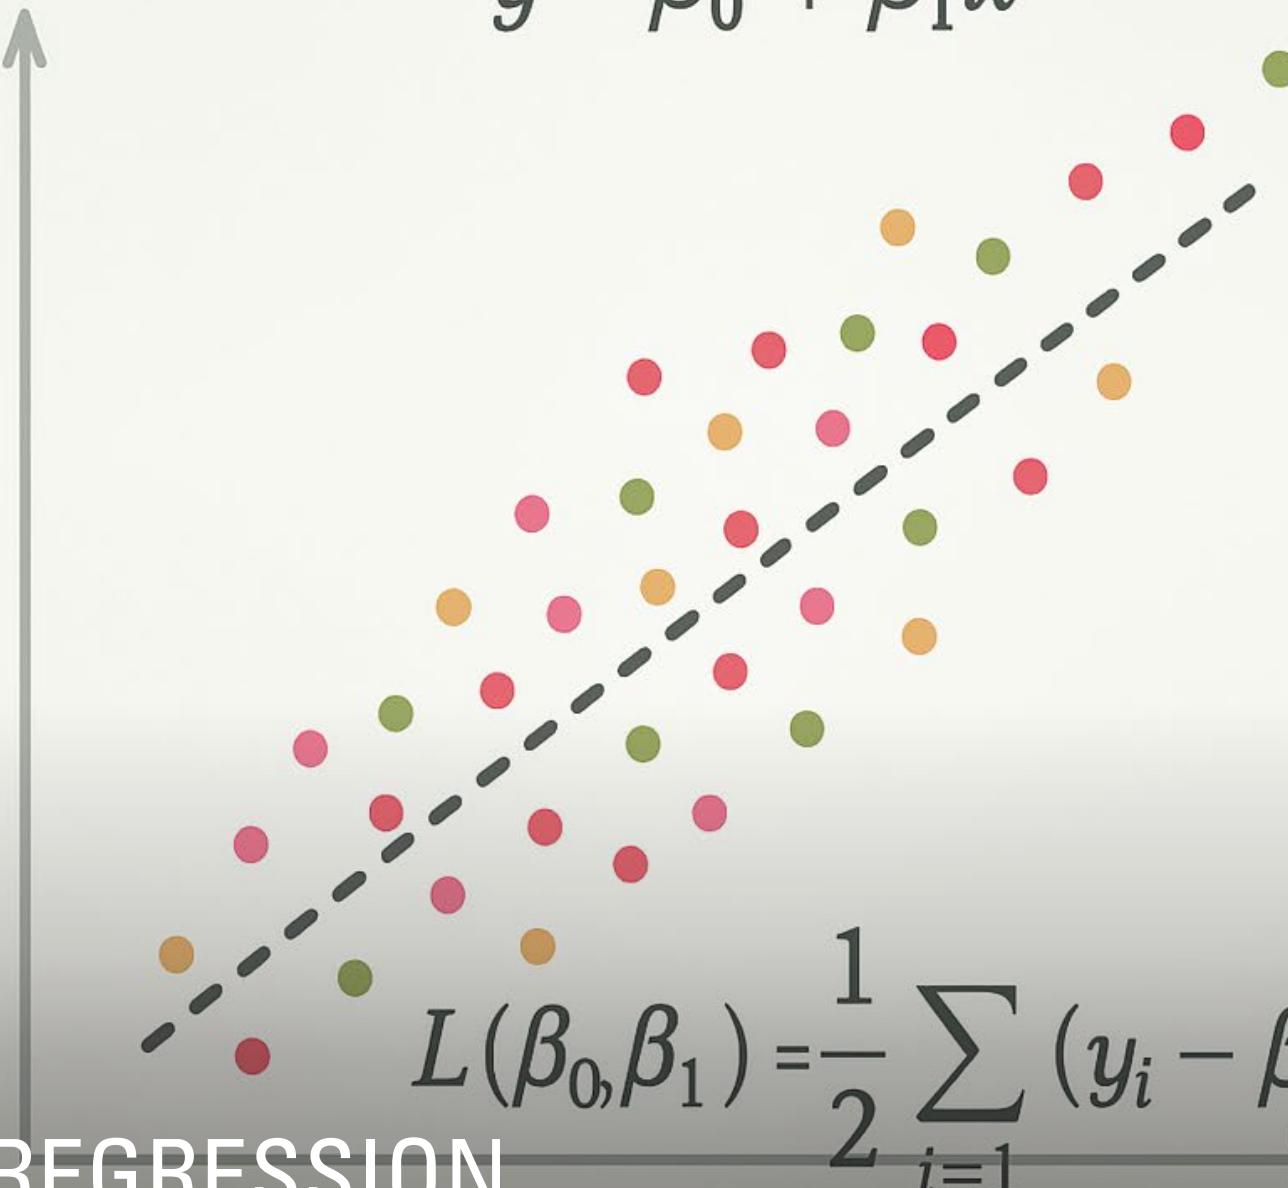


$$y = \beta_0 + \beta_1 x$$



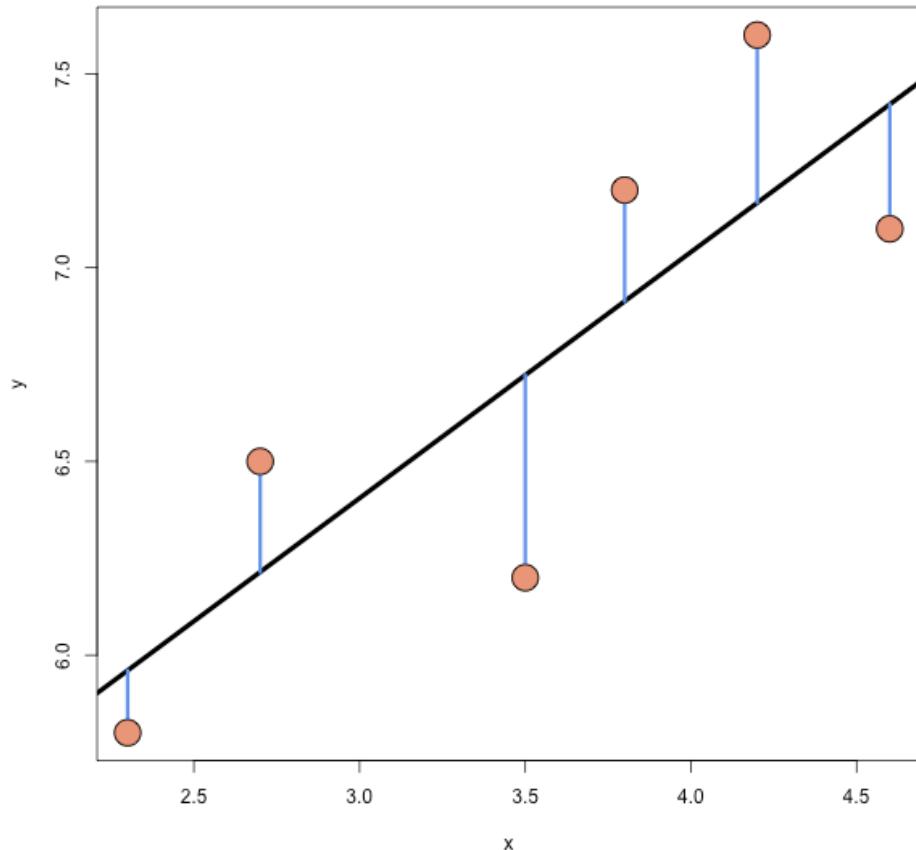
INTRO TO REGRESSION

Linear Regression

LINEAR REGRESSION

- ▶ Linear regression is a **supervised** learning approach for **predicting** a **quantitative** response
 - ▶ Linear regression has been around for a long time
 - Still very widely used
 - Best method for truly linear relationships
 - Many "fancier" methods are based on linear regression
-

LEAST SQUARES



Least squares line is the line that minimizes

$$\frac{1}{2n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Training a Linear Regression Model

GENERAL MODEL-BASED TRAINING

- Define an appropriate **objective function**
 - Cost/Loss functions are minimized (machine learning)
 - Utility/Fitness functions are maximized (statistics)
 - For example, the likelihood in MLE is a utility function
- “Train” the model by solving the optimization problem using an appropriate optimization **method**
 - a “closed-form” solution (rare)
 - a numerical method (usually)
 - **Gradient descent or variants (Adam, Newton’s method)**

OPTIMIZING THE OBJECTIVE FUNCTION

- Minimum vs. Minimizer
- Given a function, f , how do we minimize the function?
- What if there is no analytic solution? (I.e., no closed form?)
- A minimizer is the value x that produces the smallest output $f(x)$.
- Differential calculus!
- Numerical approaches.

ORDINARY LEAST SQUARES OBJECTIVE

- What does it mean to “fit” a line to data?
- How do we measure goodness?
- Why use $(y - \hat{y})^2$ and not $|y - \hat{y}|$?
- The objective function:
 - We want $\hat{y} = X\beta$ to be a *good* approximation.
 - We want $y - \hat{y}$ to be *small*.
 - Differentiability/punishing larger errors.

$$J(\beta) = \frac{1}{2n} (y - X\beta)^T (y - X\beta)$$

METHOD 1: NORMAL EQUATIONS

- Matrix notation cost function:

$$J(\beta) = \frac{1}{2n} (y - X\beta)^T (y - X\beta)$$

- Derivative of Cost Function:

$$\nabla J(\beta) = \frac{1}{n} X^T (X\beta - y)$$

- Set equal to 0 and rearrange (normal equations):

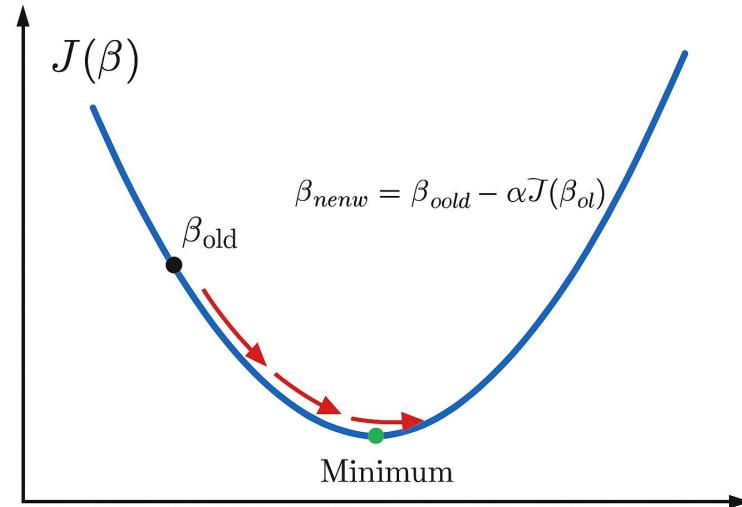
$$(X^T X)\beta = X^T y$$

- Solution for $\hat{\beta}$:

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

METHOD 2: GRADIENT DESCENT

- Numerical optimization algorithm that tweaks parameters iteratively until the minimum is found
 - “Descend” the gradient of the cost function in the steepest direction until the minimum is reached
- What is the gradient?
 - The vector that points in the direction of the greatest rate of increase for a function
 - It is the vector of partial derivatives



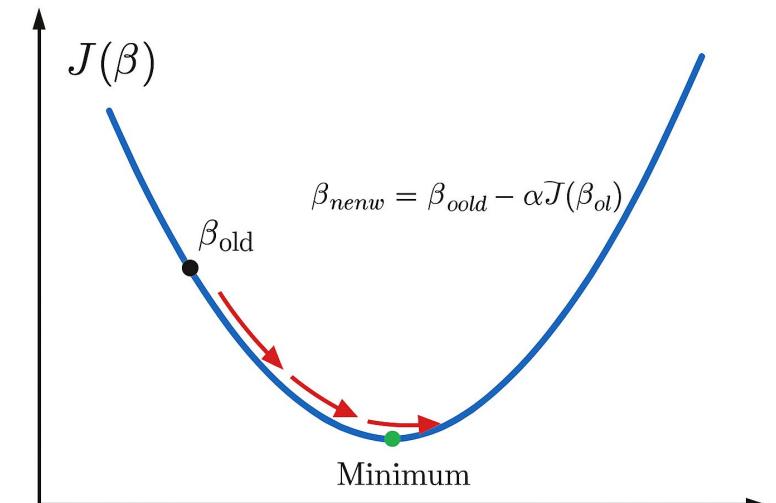
METHOD 2: GRADIENT DESCENT

- **Batch Gradient Descent Algorithm:**

- Choose initial values for the β vector (that is, choose a value for each β_j)
- Adjust β by moving in the opposite direction of the gradient:

$$\beta_{new} = \beta_{old} - \alpha \nabla J(\beta_{old})$$

- Where α is the *learning rate*, a hyperparameter, that controls the step size
- Repeat previous step until convergence is reached
- The term “Batch” refers to the fact that all the training data is used



LEARNING RATE

- If the learning rate is too high, the gradient descent algorithm will bounce around the minimum and never converge
 - (See Figure 4-5 in HOML)
 - If the learning rate is too low, it could take a very long time to find the minimum
 - (See Figure 4-4 in HOML)
 - It is common to start with a small learning rate (~ 0.01) and adjust based on the performance of the model
-

GRADIENT DESCENT VARIATIONS

- Stochastic Gradient Descent (SGD)
 - How it works:
 - Chooses one random instance at each step.
 - Computes the gradient based on that single instance.
 - **Pros:**
 - Faster per update (small computation).
 - Randomness helps escape local minima on complex cost surfaces.
 - **Cons:**
 - Updates are noisy; the algorithm never fully “settles.”
 - True minimum is rarely reached—oscillates around it.

GRADIENT DESCENT VARIATIONS

- Mini-Batch Gradient Descent
 - How it works:
 - Chooses a small random subset (mini-batch) of training data at each step.
 - Computes the gradient based on that subset.
 - **Pros:**
 - More stable than SGD (less noisy updates).
 - Faster than full batch gradient descent for large datasets.
 - Leverages vectorized operations for efficiency.
 - **Cons:**
 - Still introduces some randomness, so convergence is not perfectly smooth.
 - Requires tuning batch size (too small = noisy, too large = slow).

GRADIENT DESCENT VARIATIONS

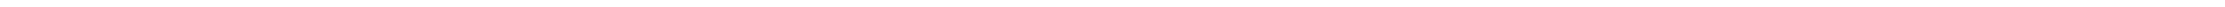
- Adam Optimizer
- How it works:
 - Maintains two moving averages:
 - First moment (m): average of gradients (like momentum).
 - Second moment (v): average of squared gradients
 - Applies bias correction to both moments.
 - Update rule:

$$\beta_{new} = \beta_{old} - \alpha \frac{\hat{m}}{\sqrt{\hat{v}} + \epsilon}$$

LEARNING SCHEDULE

- A way to help “settle” stochastic gradient descent
- Gradually reduce the learning rate according to a function called the “learning schedule”
- Additional Points:
 - Common schedules include:
 - Step decay: Reduce learning rate by a factor after a fixed number of epochs.
 - Exponential decay: Multiply learning rate by a constant factor each epoch.
 - Polynomial decay: Decrease learning rate following a polynomial function.
- Modern techniques:
 - Cosine Annealing: Learning rate follows a cosine curve, often combined with restarts.
 - One-cycle policy: Increase learning rate initially, then decrease sharply for better convergence.
 - Adaptive methods (Adam, RMSProp): Adjust learning rate per parameter automatically based on gradient history.

REGULARIZATION



REGULARIZATION

- Adding constraints to model to "simplify" the model to help prevent overfitting
- **Regularization:** Techniques that constrain a model to reduce overfitting and improve generalization.
 - Explicit Regularization
 - Implicit Regularization

EXPLICIT REGULARIZATION

- Intentionally and directly add a regularization term to the loss function
 - Explicitly modify the optimization objective
 - Examples (not exhaustive, but these are the most common):
 - L1 Regularization (Lasso)
 - L2 Regularization (Ridge)
 - Elastic Net (combine L1 and L2 regularization)
-

IMPLICIT REGULARIZATION

- Any regularization method that does not explicitly modify the cost or objective function
- Can be a byproduct of an algorithm, the learning process, or other choices made in training
- Examples (not exhaustive):
 - Early stopping (in gradient-based optimization, stop training before convergence is reached)
 - Dropout (randomly ignoring some parts of a neural network)
 - Pruning (simplifying a grown decision tree)
 - Data augmentation (adding synthetic data to increase the sample size)

Example of Explicit Regularization in Linear Regression

REGULARIZED COST FUNCTIONS

- Least squares cost function: finds the β 's that minimize

$$J(\beta) = \frac{1}{n} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2$$

L1 PENALTY: LASSO

- Least squares cost function

$$J(\beta) = \frac{1}{n} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2$$

Tuning parameter that controls relative impact of two criteria

- Lasso linear regression cost function

$$J_{lasso}(\beta) = \frac{1}{n} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \alpha \sum_{j=1}^p |\beta_j|, \quad (\alpha > 0)$$

Tries to fit the data well
(find β 's that move the prediction close to the observed data)

What happens if $\alpha = 0$ or $\alpha \approx \infty$?

Constrains the model
(pushes β 's to be close to or exactly equal to 0)

REGULARIZED COST FUNCTIONS

ℓ_1 penalty

- Lasso linear regression cost function

$$J_{lasso}(\beta) = \frac{1}{n} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \alpha \sum_{j=1}^p |\beta_j|, \text{ (where } \alpha > 0\text{)}$$

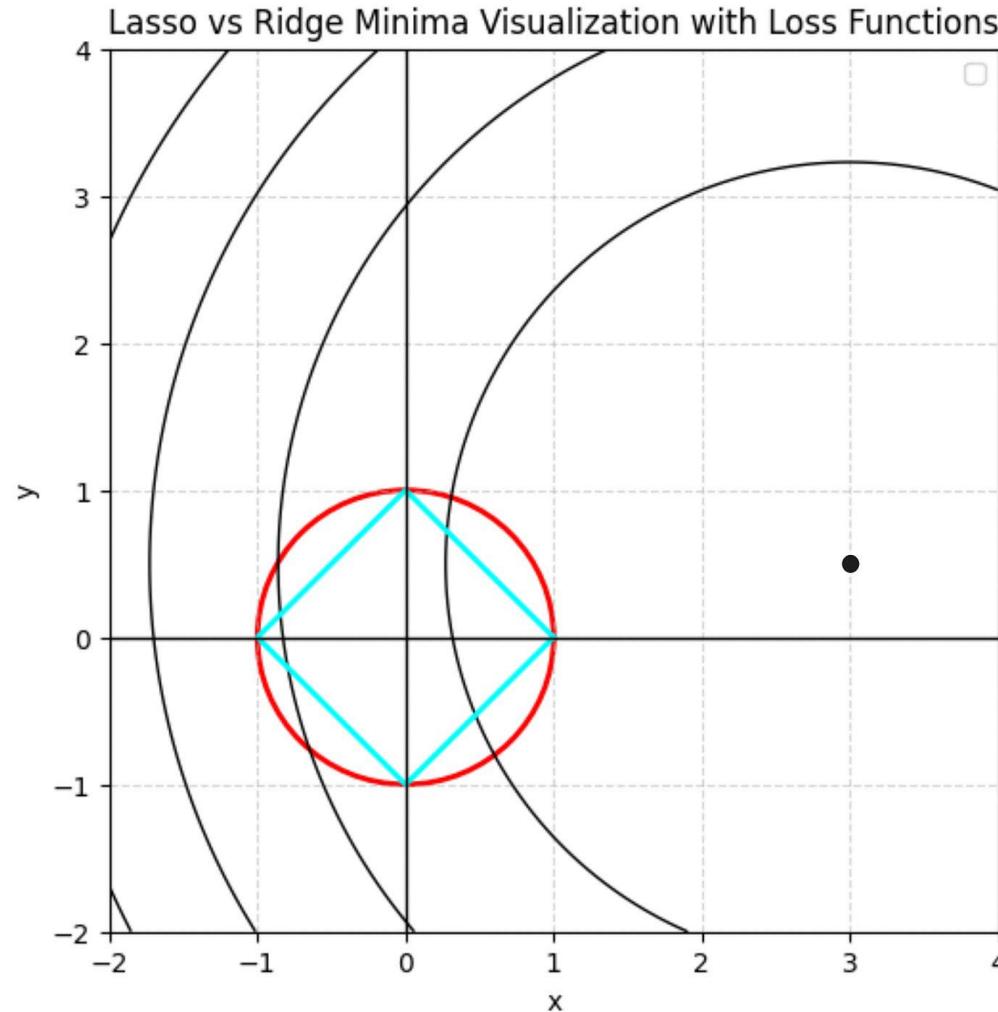
ℓ_2 penalty

- Ridge linear regression cost function

$$J_{ridge}(\beta) = \frac{1}{n} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \alpha \sum_{j=1}^p (\beta_j)^2, \text{ (where } \alpha > 0\text{)}$$


Constrains the model
(pushes β s to be close to but NOT
exactly equal to 0)

LASSO V. RIDGE VISUALIZATION



Let's explore this further in
Desmos:
<https://www.desmos.com/3d>

<https://www.vizuaranewsletter.com/p/it-regularization-is-easy-can-you>: This is a great source for an intuitive description!

RIDGE VS LASSO

- Lasso
 - Forces some coefficients to be exactly 0 (advantage)
 - Built in variable selection
 - Coefficients can be very different for correlated features (disadvantage)
 - Ridge
 - Coefficients will never be exactly 0 (fact)
 - Tends to estimate similar coefficient for correlated features (advantage)
 - Neither method will universally dominate
 - Lasso generally performs better when many of the features don't contribute to predictability
 - Ridge generally performs better when many of the features contribute to the prediction and when features are correlated
-

ELASTIC-NET REGULARIZATION

- Combination of Ridge and Lasso
 - The idea is to capitalize on the strengths of both methods

$$J_{en}(\beta) = \frac{1}{n} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \alpha_1 \sum_{j=1}^p |\beta_j| + \alpha_2 \sum_{j=1}^p (\beta_j)^2$$

- If $\alpha_1 = 0$ (and $\alpha_2 \neq 0$), elastic net becomes Ridge
- If $\alpha_2 = 0$ (and $\alpha_1 \neq 0$), elastic net becomes Lasso
- If both are 0, we're back to vanilla least squares

REGULARIZATION GENERALLY

- We've just looked at Lasso, Ridge, and Elastic net (explicit regularization) in the context of linear regression
- Regularization is a general set of methodologies, not limited to linear regression
 - Used in many models: logistic regression, neural networks, decision trees, etc.
 - Goal: improve generalization by reducing overfitting
 - Can be applied to weights, complexity, or even the training process

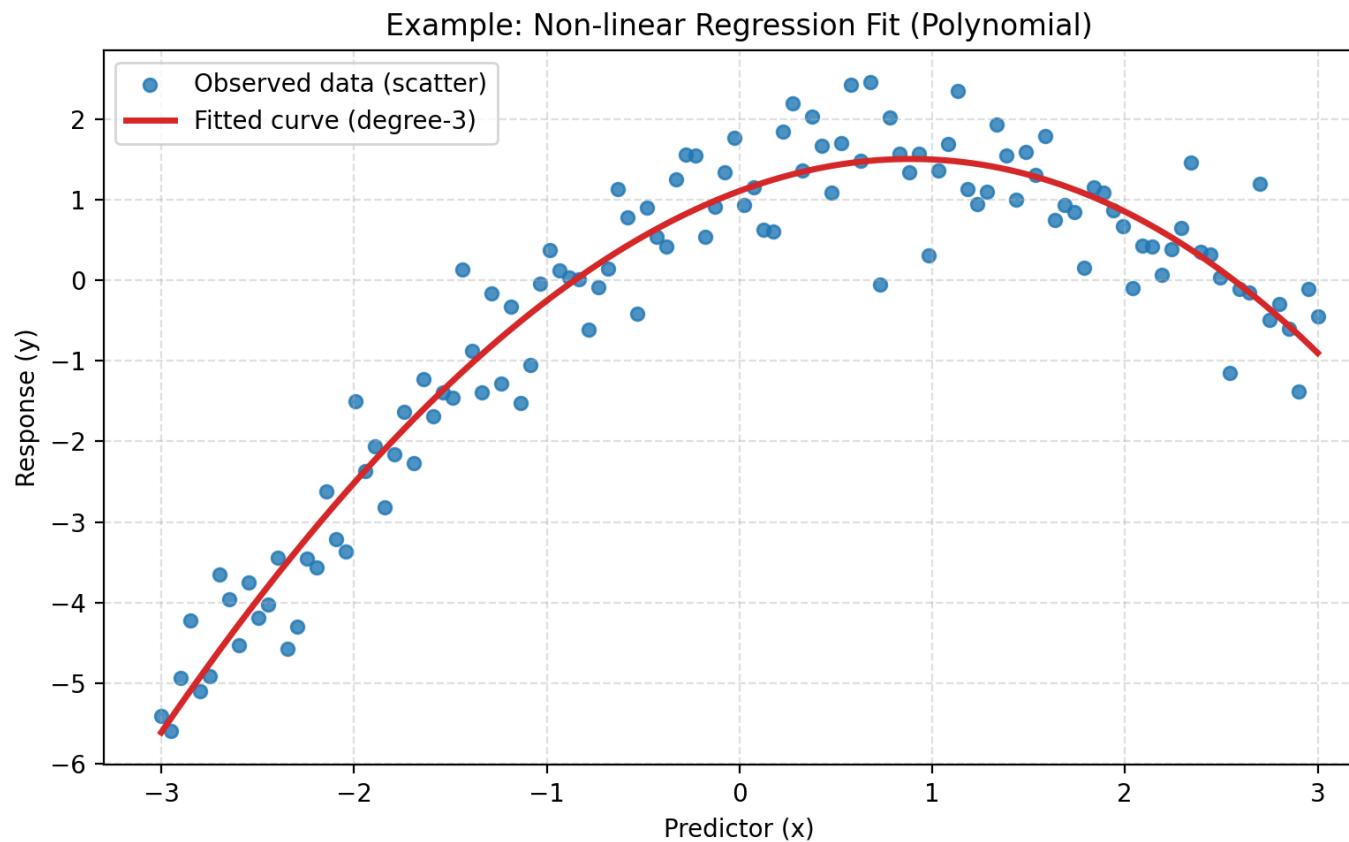
“NONLINEAR” LINEAR REGRESSION

POLYNOMIAL REGRESSION

- Sometimes the relationship between X and y is non-linear
- We can use linear regression to model non-linear relationships (between X and y)
 - Add quadratic, cubic, etc. terms to the model
 - This is sometimes called polynomial regression
 - (Now we are moving into the realm of multiple regression, where we have more than one predictor)

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \beta_3 x_i^3 + \epsilon_i$$

POLYNOMIAL REGRESSION



$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \epsilon_i$$

TERMINOLOGY

- ▶ Linear regression is "linear" with respect to the parameters (the model is a linear combination of the parameters)

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \beta_3 x_i^3 + \epsilon_i$$

- ▶ True "nonlinear" regression is where the regression model is a nonlinear combination of the parameters, for example:

$$y_i = \frac{\beta_1 x_i}{\beta_0 + x_i} + \epsilon_i$$

MORE TERMINOLOGY

- ▶ Just FYI:
 - **Multiple** regression is a regression model with more than one predictor (or feature)
 - **Multivariate** regression is a regression model with more than one response (or target)

INTERACTION & POLYNOMIAL FEATURES

```
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures

# Add squared and cubed features for every predictor plus interactions
poly = PolynomialFeatures(degree=3)
X = poly.fit_transform(X)

# Only include up to 3rd-order interactions
inter = PolynomialFeatures(degree=3, interaction_only=True)
X = inter.fit_transform(X)
```

REGRESSION METRICS

EVALUATING A REGRESSION MODEL

- In regression (we are predicting a numeric variable),
what does a “good” prediction look like?
- Ideally, \hat{y} , (the prediction) is close to y (the observed value)

COMMON REGRESSION METRICS

- Lower is better:

- Mean absolute error (MAE): $MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$

- Mean squared error (MSE): $MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$

- Root mean squared error (RMSE): $RMSE = \sqrt{MSE}$

- Higher is better:

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

REGRESSION METRIC COMPARISONS

MAE (Mean Absolute Error)

- ✓ Easy to interpret (average absolute difference)
- ✓ Less sensitive to outliers than MSE
- ✗ Does not penalize large errors as strongly

R² (Coefficient of Determination)

- ✓ Measures proportion of variance explained
- ✓ Easy to compare across models
- ✗ Can be misleading for non-linear

MSE (Mean Squared Error)

- ✓ Penalizes large errors more (squared term)
- ✓ Commonly used in optimization (differentiable)
- ✗ Sensitive to outliers

RMSE (Root Mean Squared Error)

- ✓ Same units as target variable
- ✓ Highlights large errors clearly
- ✗ Still sensitive to outliers

Regression metrics

User guide. See the [Regression metrics](#) section for further details.

d2_absolute_error_score	D^2 regression score function, fraction of absolute error explained.
d2_pinball_score	D^2 regression score function, fraction of pinball loss explained.
d2_tweedie_score	D^2 regression score function, fraction of Tweedie deviance explained.
explained_variance_score	Explained variance regression score function.
max_error	The max_error metric calculates the maximum residual error.
mean_absolute_error	Mean absolute error regression loss.
mean_absolute_percentage_error	Mean absolute percentage error (MAPE) regression loss.
mean_gamma_deviance	Mean Gamma deviance regression loss.
mean_pinball_loss	Pinball loss for quantile regression.
mean_poisson_deviance	Mean Poisson deviance regression loss.
mean_squared_error	Mean squared error regression loss.
mean_squared_log_error	Mean squared logarithmic error regression loss.
mean_tweedie_deviance	Mean Tweedie deviance regression loss.
median_absolute_error	Median absolute error regression loss.
r2_score	R^2 (coefficient of determination) regression score function.
root_mean_squared_error	Root mean squared error regression loss.
root_mean_squared_log_error	Root mean squared logarithmic error regression loss.

REGRESSION METRICS IN SCIKIT-LEARN

In v1.4 or greater
(previous versions
have a “squared”
argument in the
MSE function)



MODEL SELECTION METRICS

- Models selection tools, such as “GridSearchCV” and “cross_val_score” use a scoring parameter

For the most common use cases, you can designate a scorer object with the `scoring` parameter; the table below shows all possible values. All scorer objects follow the convention that **higher return values are better than lower return values**. Thus metrics which measure the distance between the model and the data, like `metrics.mean_squared_error`, are available as `neg_mean_squared_error` which return the negated value of the metric.

'neg_mean_absolute_error'	<code>metrics.mean_absolute_error</code>
'neg_mean_squared_error'	<code>metrics.mean_squared_error</code>
'neg_root_mean_squared_error'	<code>metrics.root_mean_squared_error</code>